

# Map Reduce

---

## Table of Contents

- **Map Reduce**
  - **1. Map Reduce for Word Count**
    - **1.1 Describe what might have caused this issue.**
    - **1.2 Suggest a possible solution to this problem. You need to specifically look in to the distributed system features.**
    - **1.3 Comment on the efficiency of the database that is utilized here and suggest a better approach to handle it.**

# 1. Map Reduce for Word Count

Assume you are tasked to write a MapReduce-based application (a Map function and a Reduce function) that analyzes the way words are used in documents. Each input document corresponds to a separate Map task where Reduce function needs to know the total number of times each word occurs in the total set of documents.

A database server is setup that has a table mapping words to counts. Map function is programmed to connect to that database over the network and tell it to increment the relevant count in the database for each word that Map sees in its input. Reduce function is written to consult the database when it needs a word's total count. Counts are reset carefully in the database to zero before starting a MapReduce job. A single MapReduce job is run at a time.

However, to your surprise, you find that sometimes the counts that the Reduces read from the database are too high – higher than the total number of times a word actually appears.

## 1.1 Describe what might have caused this issue.

Assuming only a single instance of the 'Map' is running at a time, that means whenever the 'Map' finds the word, it triggers an 'IncDB(word)' i.e. Only one increment at a time. If there's a network failure, the message will not be recieved by the DB, which means no increment takes place. That means, as long as there's only 1 'Map' task running, the DB is safe from incorrectly writing a value greater than the actual value to the DB, because a single 'Map' cannot incorrectly trigger an increment more than once, whenever it finds the matching word.

However, even if the issue is about reading 'higher than the total number of times a word actually appears', with a single 'Map' task, if sent 'IncDB(word)' fails, it means that the actual increment is not written to the DB i.e. Reduce reads lower than the total numner of times a word actually appears.

msg	word	recieved	count
-	w	-	0 (init)
1, IncDB(w)	w	t	1
2, IncDB(w)	w	t	2
3, IncDB(w)	w	t	3
4, IncDB(w)	w	f	3
5, IncDB(w)	w	t	4

So, the issue has to occur due to running multiple 'Map' tasks even though it is said to be 'single instance at a time'.

## 1.2 Suggest a possible solution to this problem. You need to specifically look in to the distributed system features.

The 'Map' task only triggers the 'IncDB(word)' function whenever the word is found. Assuming there are 3 'Map' tasks running at them same time but with some transmission delays in delivering the message to DB.

Map task	delay
M1	0.5 s
M2	0.0 s
M3	0.8 s

Now to fix the issue, we can have the 'IncDB(word)' to keep track of the last read value in the DB i.e convert 'IncDB(word)' into 'IncDB(lastCount, word)'

This way, the increment can be prevented if the lastCount value is not greater than or equal to the actual DB value i.e. prevents 'Map' tasks from writing higher counts than the actual.

It's important to note that LOCK/UNLOCK alone cannot solve this issue i.e. This would only prevent the multiple 'Map' tasks to access the specific record simultaneously. It doesn't prevent the other 'Map' tasks to incorrectly increment the value eventually.

Another way to solve this is to make sure the 'Master' kills any 'Workers' who run the same 'Map' task i.e. The 'Master' pings the 'Workers' periodically and whenever 'Master' finds a 'Worker' who runs the same 'Map' task, that 'Worker' must be immediately killed. However, this still doesn't solve the issue for lost messages that causes the 'Reducer' to read low count than the actual count.

## 1.3 Comment on the efficiency of the database that is utilized here and suggest a better approach to handle it.

Based on the problem scenario, we can assess which factors are likely available and which are not available in the given database:

Available:

1. Concurrency: The database allows multiple Map tasks to access and update word counts concurrently over the network.
2. Network Latency
3. Data Model: The database has a schema mapping words to counts

Not available:

1. Scalability: does not explicitly mention how to accommodate a large number of concurrent requests.
2. Locking Mechanisms: no mention of specific locking mechanisms.
3. Indexing and Query Optimization: There's no indication of optimization strategies for queries or indexing mechanisms within the database.
4. Fault Tolerance: There's an indication of handling failures by resetting counts to zero before each MapReduce job, suggesting some level of fault tolerance.

So a better database with the following properties would be ideal to tackle the efficiency drawback in the existing database.

A distributed key-value store with strong concurrency controls to manage simultaneous updates from Map tasks. Implement effective locking to maintain consistency and prevent overcounting. Making sure the database is horizontally scalable to handle a high volume of concurrent requests and data. Ensure fault tolerance with features like automatic failover and data replication. Optimize indexing and query performance for efficient word count retrieval. Implement monitoring and management tools for proactive maintenance. This database design will efficiently support word count processing in the MapReduce job while ensuring scalability, fault tolerance, and data consistency.